

L07 Tidy Data

Data Science I (STAT 301-1)

Contents

Overview	1
Datasets	1
Exercises	1

Overview

The goal of this lab is to learn what it means to be a “tidy” dataset and how to tidy messy datasets utilizing the `tidyr` package — a core member of the `tidyverse`. [tidyr package home page](#)

Datasets

All datasets are either defined inline or provided within the core `tidyverse` packages (`table1`, `table2`, `table4a`, `table4a`, `who`)

Exercises

Please complete the following exercises. Be sure your solutions are clearly indicated and that the document is neatly formatted.

```
library(tidyverse)
```

Load Packages

Exercise 1 (Website: 12.2.1 Ex. 2 — modified)

Follow these four steps to compute the `rate` per 10,000 once using only `table2` and again using `table4a` + `table4b`:

1. Extract the number of TB cases per country per year.
2. Extract the matching population numbers per country per year.
3. Divide case numbers by population numbers and multiply by 10000.
4. Store in a new `tibble`.

How does this process compare to using `table1`, which is a tidy dataset, to compute `rate` per 10,000?

```
#first using table2
case_per_year <- table2 %>%
  filter(type == 'cases') %>%
```

```

group_by(year, country) %>%
  summarise(total_cases = sum(count))

pop_per_year <- table2 %>%
  filter(type == 'population') %>%
  group_by(year, country) %>%
  summarise(total_pop = sum(count))

rate <- full_join(case_per_year, pop_per_year)
rate %>%
  mutate(rate = 10000*total_cases/total_pop)

## # A tibble: 6 x 5
## # Groups:   year [2]
##   year country    total_cases total_pop  rate
##   <int> <chr>         <int>      <int> <dbl>
## 1  1999 Afghanistan      745    19987071 0.373
## 2  1999 Brazil          37737   172006362 2.19
## 3  1999 China          212258  1272915272 1.67
## 4  2000 Afghanistan      2666    20595360 1.29
## 5  2000 Brazil          80488   174504898 4.61
## 6  2000 China          213766  1280428583 1.67

```

Next using table4a and table4b

```

tidy4a <- table4a %>%
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "cases")

tidy4b <- table4b %>%
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "population")

rate2 <- left_join(tidy4a, tidy4b)

## Joining, by = c("country", "year")

rate2 %>%
  mutate(rate = 10000*cases/population)

```

```

## # A tibble: 6 x 5
##   country    year  cases population  rate
##   <chr>      <chr> <int>      <int> <dbl>
## 1 Afghanistan 1999     745    19987071 0.373
## 2 Afghanistan 2000     2666    20595360 1.29
## 3 Brazil       1999    37737   172006362 2.19
## 4 Brazil       2000    80488   174504898 4.61
## 5 China        1999   212258  1272915272 1.67
## 6 China        2000   213766  1280428583 1.67


```

Using table1 would be a lot easier, as you wouldn't have to filter by type.


Exercise 2 (Website: 12.2.1 Ex. 3)

Recreate the plot below showing the change in cases over time using table2 instead of table1. What do you need to do first?

```
# Change over time in number of TB cases by country
ggplot(table1, aes(year, cases)) +
  geom_line(aes(group = country), colour = "grey50") +
  geom_point(aes(colour = country))
```



```
table2 %>%
  filter(type == 'cases') %>%
  ggplot(aes(year, count)) +
  geom_line(aes(group = country), colour = 'grey50') +
  geom_point(aes(colour = country))
```



First you need to filter so that only the cases are plotted, not the populations.

Exercise 3 (Website: 12.3.3 Ex. 2)

Why does the provided code fail? Fix it.

```
table4a %>%
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "cases")
```

```
## # A tibble: 6 x 3
##   country    year  cases
##   <chr>      <chr> <int>
## 1 Afghanistan 1999     745
## 2 Afghanistan 2000    2666
## 3 Brazil      1999   37737
## 4 Brazil      2000   80488
## 5 China       1999  212258
## 6 China       2000  213766
```

You need backticks on the 1999 and 2000, or else R doesn't recognize them as variable names.

Exercise 4 (Website: 12.3.3 Ex. 3)

What would happen if you use `pivot_wider()` on this table? Add a new column to fix the problem.

```
people <- tribble(
  ~name,      ~key,    ~value,
  #-----/-----/-----
```

```

"Phillip Woods", "age", 45,
"Phillip Woods", "height", 186,
"Phillip Woods", "age", 50,
"Jessica Cordero", "age", 37,
"Jessica Cordero", "height", 156
)

```

```

people %>%
  pivot_wider(names_from = key, values_from = value)

```

```

## Warning: Values are not uniquely identified; output will contain list-cols.
## * Use `values_fn = list` to suppress this warning.
## * Use `values_fn = length` to identify where the duplicates arise
## * Use `values_fn = {summary_fun}` to summarise duplicates

```

```

## # A tibble: 2 x 3
##   name          age      height
##   <chr>         <list>   <list>
## 1 Phillip Woods <dbl [2]> <dbl [1]>
## 2 Jessica Cordero <dbl [1]> <dbl [1]>

```

I believe it is not working because Phillip Woods has two values for age, and R doesn't know where to put the second value. I will make a variable called obs that accounts for the two ages of Phillip Woods.

```

people2 <- people %>%
  group_by(name, key) %>%
  mutate(obs = row_number())

```

```

people2 %>%
  pivot_wider(names_from = "name", values_from = "value")

```

```

## # A tibble: 3 x 4
## # Groups:   key [2]
##   key      obs `Phillip Woods` `Jessica Cordero`
##   <chr> <int>         <dbl>         <dbl>
## 1 age      1           45           37
## 2 height   1          186          156
## 3 age      2           50           NA

```

This table still doesn't look great despite the fix.

Exercise 5 (Website: 12.3.3 Ex. 4)

Tidy the simple tibble below. Do you need to make it wider or longer? What are the variables in your tidy version?

```

preg <- tribble(
  ~pregnant, ~male, ~female,
  "yes",      NA,    10,
  "no",       20,    12
)

```

```

preg %>%
  pivot_longer(c(male, female), names_to = 'Gender', values_to = 'Count')

```

```
## # A tibble: 4 x 3
##   pregnant Gender Count
##   <chr>    <chr> <dbl>
## 1 yes     male     NA
## 2 yes     female   10
## 3 no      male    20
## 4 no      female   12
```

You need to make it longer. I made new variables called Gender and Count.

Exercise 6 (Website: 12.4.3 Ex. 1)

What do the extra and fill arguments do in `separate()`? Experiment with the various options for the following two datasets.

```
tibble(x = c("a,b,c", "d,e,f,g", "h,i,j")) %>%
  separate(x, c("one", "two", "three"), extra = 'warn')
```

```
## Warning: Expected 3 pieces. Additional pieces discarded in 1 rows [2].
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    f
## 3 h    i    j
```

```
tibble(x = c("a,b,c", "d,e,f,g", "h,i,j")) %>%
  separate(x, c("one", "two", "three"), extra = 'drop')
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    f
## 3 h    i    j
```

```
tibble(x = c("a,b,c", "d,e,f,g", "h,i,j")) %>%
  separate(x, c("one", "two", "three"), extra = 'merge')
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    f,g
## 3 h    i    j
```

```
tibble(x = c("a,b,c", "d,e", "f,g,i")) %>%
  separate(x, c("one", "two", "three"), fill = 'warn')
```

```
## Warning: Expected 3 pieces. Missing pieces filled with `NA` in 1 rows [2].
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    <NA>
## 3 f    g    i
```

```
tibble(x = c("a,b,c", "d,e", "f,g,i")) %>%
  separate(x, c("one", "two", "three"), fill = 'right')
```

```
## # A tibble: 3 x 3
##   one    two    three
##   <chr> <chr> <chr>
## 1 a      b      c
## 2 d      e      <NA>
## 3 f      g      i
```

```
tibble(x = c("a,b,c", "d,e", "f,g,i")) %>%
  separate(x, c("one", "two", "three"), fill = 'left')
```

```
## # A tibble: 3 x 3
##   one    two    three
##   <chr> <chr> <chr>
## 1 a      b      c
## 2 <NA>   d      e
## 3 f      g      i
```

extra controls what happens if there are too many pieces. Setting it to 'warn' (the default), shows the warning and drops extra values, to 'drop' drops extra values without a warning, and to 'merge' only splits at most length(into) times. Note how 'warn' and 'drop' omitted the 'g' but merge added it. fill is similar, except it controls what happens when there aren't enough pieces. Setting it to 'warn' is the default, which will fill from the right. You can also set it to 'right' and 'left'.

Exercise 7 (Website: 12.4.3 Ex. 2)

Both unite() and separate() have a remove argument. What does it do? Why would you set it to FALSE? remove removes the input column from the output data frame. If you don't want to lose that column, set it to FALSE.

Exercise 8 (Website: 12.6)

The case study of data from the 2014 World Health Organization Global Tuberculosis Report produces a lot of useful data (who) — data sub-directory contains the codebook for who. However, the format is difficult to work with, so the authors walk you through the process of tidying the data. The tidying process uses this nice concise code:

```
who %>%
  gather(code, value, new_sp_m014:newrel_f65, na.rm = TRUE) %>%
  mutate(code = stringr::str_replace(code, "newrel", "new_rel")) %>%
  separate(code, c("new", "var", "sexage")) %>%
  select(-new, -iso2, -iso3) %>%
  separate(sexage, c("sex", "age"), sep = 1)
```

```
## # A tibble: 76,046 x 6
##   country    year var    sex    age    value
##   <chr>      <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp    m     014      0
## 2 Afghanistan 1998 sp    m     014     30
## 3 Afghanistan 1999 sp    m     014      8
## 4 Afghanistan 2000 sp    m     014     52
```

```
## 5 Afghanistan 2001 sp m 014 129
## 6 Afghanistan 2002 sp m 014 90
## 7 Afghanistan 2003 sp m 014 127
## 8 Afghanistan 2004 sp m 014 139
## 9 Afghanistan 2005 sp m 014 151
## 10 Afghanistan 2006 sp m 014 193
## # ... with 76,036 more rows
```

Insert a comment following each # in the code below that explains the purpose or objective of the line of code directly below it.

```
who_tidy <- who %>%
  # condenses the columns from `new_sp_m014` to `newrel_f65` into two new variables called `code` and
  pivot_longer(cols = c(new_sp_m014:newrel_f65),
    names_to = "code",
    values_to = "value",
    values_drop_na = TRUE) %>%
  #replace values with 'newrel' to 'new_rel', for consistency.
  mutate(code = stringr::str_replace(code, "newrel", "new_rel")) %>%
  # separates 'code' into three new variables, 'new', 'var', and 'sexage', based on underscore separator
  separate(code, c("new", "var", "sexage")) %>%
  #delete 'new', 'iso2', and 'iso3' variables
  select(-new, -iso2, -iso3) %>%
  #Separate 'sexage' variable into 'sex' and 'age', separate in between first and second index
  separate(sexage, c("sex", "age"), sep = 1)
```

```
who_tidy
```

```
## # A tibble: 76,046 x 6
##   country      year var  sex  age  value
##   <chr>      <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp   m    014     0
## 2 Afghanistan 1997 sp   m   1524    10
## 3 Afghanistan 1997 sp   m   2534     6
## 4 Afghanistan 1997 sp   m   3544     3
## 5 Afghanistan 1997 sp   m   4554     5
## 6 Afghanistan 1997 sp   m   5564     2
## 7 Afghanistan 1997 sp   m    65     0
## 8 Afghanistan 1997 sp   f    014     5
## 9 Afghanistan 1997 sp   f   1524    38
## 10 Afghanistan 1997 sp   f   2534    36
## # ... with 76,036 more rows
```

Exercise 9 (Website: 12.6.1 Ex. 1)

In the WHO case study, the authors set `na.rm = TRUE` to make it easier to check that they had the correct values. Is this reasonable? Think about how missing values are represented in this dataset. Are there implicit missing values? What's the difference between NA and zero in this dataset? **The appropriateness of using `na.rm = TRUE` depends on what the NA's mean. Do they mean no cases of TB or just missing data? To test if it is no cases, I will search for 0's in the dataset.**

```
who_tidy %>%
  filter(value == 0)
```

```
## # A tibble: 11,080 x 6
```



```
##   country      year var  sex  age  value
##   <chr>        <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp   m    014    0
## 2 Afghanistan 1997 sp   m    65     0
## 3 Afghanistan 1997 sp   f   5564    0
## 4 Afghanistan 2007 sn   m    014    0
## 5 Afghanistan 2007 sn   m   1524    0
## 6 Afghanistan 2007 sn   m   2534    0
## 7 Afghanistan 2007 sn   m   3544    0
## 8 Afghanistan 2007 sn   m   4554    0
## 9 Afghanistan 2007 sn   m   5564    0
## 10 Afghanistan 2007 sn   m    65     0
## # ... with 11,070 more rows
```

Obviously there are zeros, which tells me that the NA's represent missing data, in which case it would be appropriate to set `na.rm = TRUE`.

```
dim(who_tidy)
```

```
## [1] 76046      6
```

```
dim(complete(who_tidy, country, year))
```

```
## [1] 80008      6
```

From this analysis, we see that completing the missing values from country and year adds rows to the dataset, indicating that there are implicit missing values.

Exercise 10 (Website: 12.6.1 Ex. 2)

In the code from the WHO case study, what happens if you neglect the `mutate()` step? (`mutate(key = stringr::str_replace(key, "newrel", "new_rel"))`)

```
who_no_mutate <- who %>%
  pivot_longer(cols = c(new_sp_m014:newrel_f65),
               names_to = "code",
               values_to = "value",
               values_drop_na = TRUE) %>%
  separate(code, c("new", "var", "sexage")) %>%
  select(-new, -iso2, -iso3) %>%
  separate(sexage, c("sex", "age"), sep = 1)
```

```
## Warning: Expected 3 pieces. Missing pieces filled with `NA` in 2580 rows [243,
## 244, 679, 680, 681, 682, 683, 684, 685, 686, 687, 688, 689, 690, 691, 692, 903,
## 904, 905, 906, ...].
```

When you don't use `mutate`, 'newrel' will not have the same format at the other codes, and thus won't be separated properly into 'new', 'var', and 'sexage'. Thus, we'll get errors like we saw in Exercise 6.

Exercise 11 (Website: 12.6.1 Ex. 3)

The authors of the WHO case study claimed that `iso2` and `iso3` were redundant with `country`. Confirm this claim.

```
who %>%
  count(country, iso2, iso3) %>%
  count(country) %>%
  filter(n > 1)
```

```
## # A tibble: 0 x 2
```

```
## # ... with 2 variables: country <chr>, n <int>
```

There are no observations with more than one distinct combination of `country`, `iso2`, and `iso3`, thus they are redundant.

Exercise 12 (Website: 12.6.1 Ex. 4)

For each level of `country`, `year`, and `sex`, compute the total number of cases of TB (using the WHO case study data). Construct an informative visualization.

```
who_tidy %>%
  group_by(country, year, sex) %>%
  summarize(n = sum(value)) %>%
  filter(n > 50000) %>%
  ggplot(mapping = aes(x = year, y = n, color = country)) +
  geom_line(alpha = 0.5) +
  facet_wrap(~sex) +
  theme(legend.position = 'bottom')
```

