

L16 Many Models

Data Science I (STAT 301-1)

Shay Lebovitz

Contents

Overview	1
Datasets	1
Exercises	1

Overview

The goal of this lab is for students to explore and develop three powerful ideas/techniques:

1. Using many simple models to better understand complex data.
2. Using list-columns to store arbitrary data structures.
3. Using the **broom** package to turn models into tidy data.

These skills and ideas will be essential as we proceed to data science work that is beyond EDA.

Datasets

We will be utilizing the **gapminder** dataset from the **gapminder** package (you may need to install this package). See `?gapminder` for details concerning the dataset.

Exercises

Please complete the following exercises. Be sure that your solutions are clearly indicated and that the document is neatly formatted.

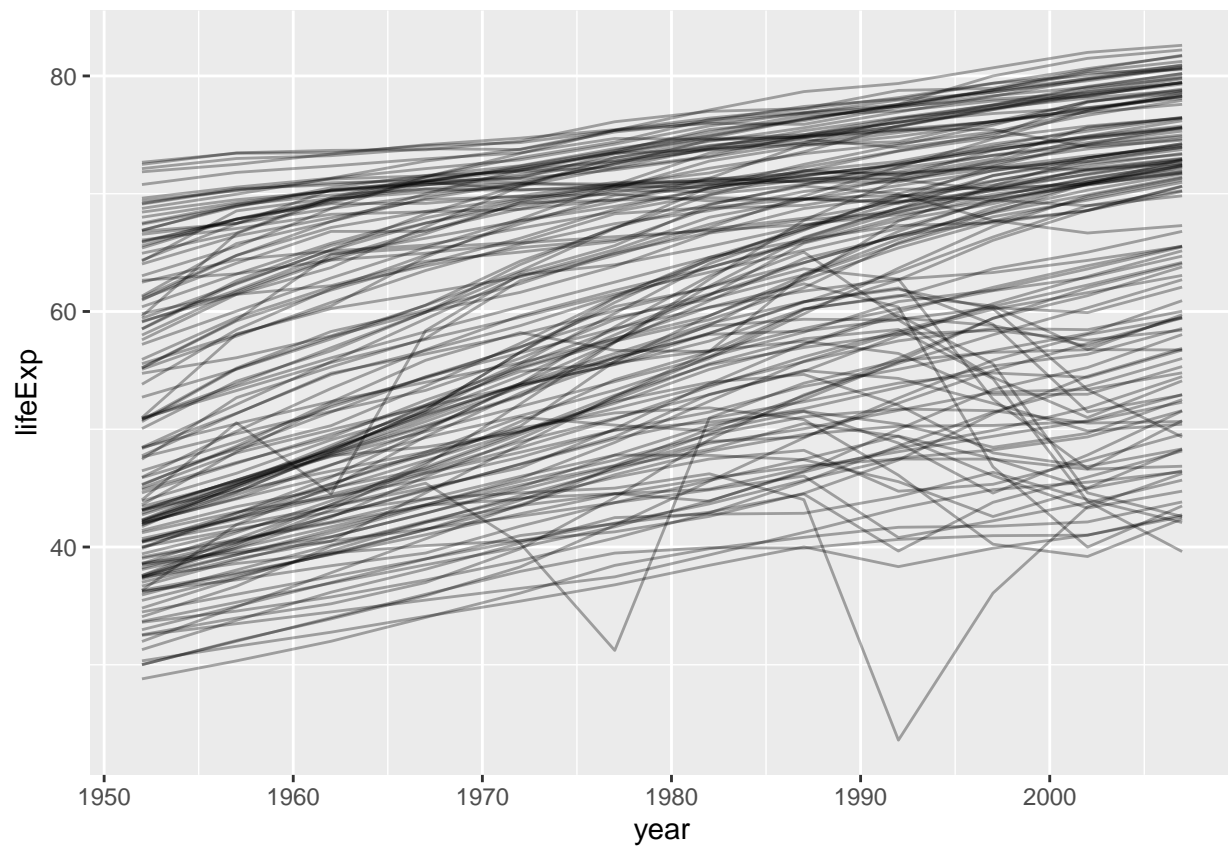
```
# Loading package(s)
library(modelr)
library(tidyverse)
library(gapminder)
library(purrr)
```

Load Packages

Exercise 1

Work through and transcribe (do not copy and paste – rewrite) the work detailed in section 25.2 gapminder. Consider adding your own comments/notes and using your own naming conventions.

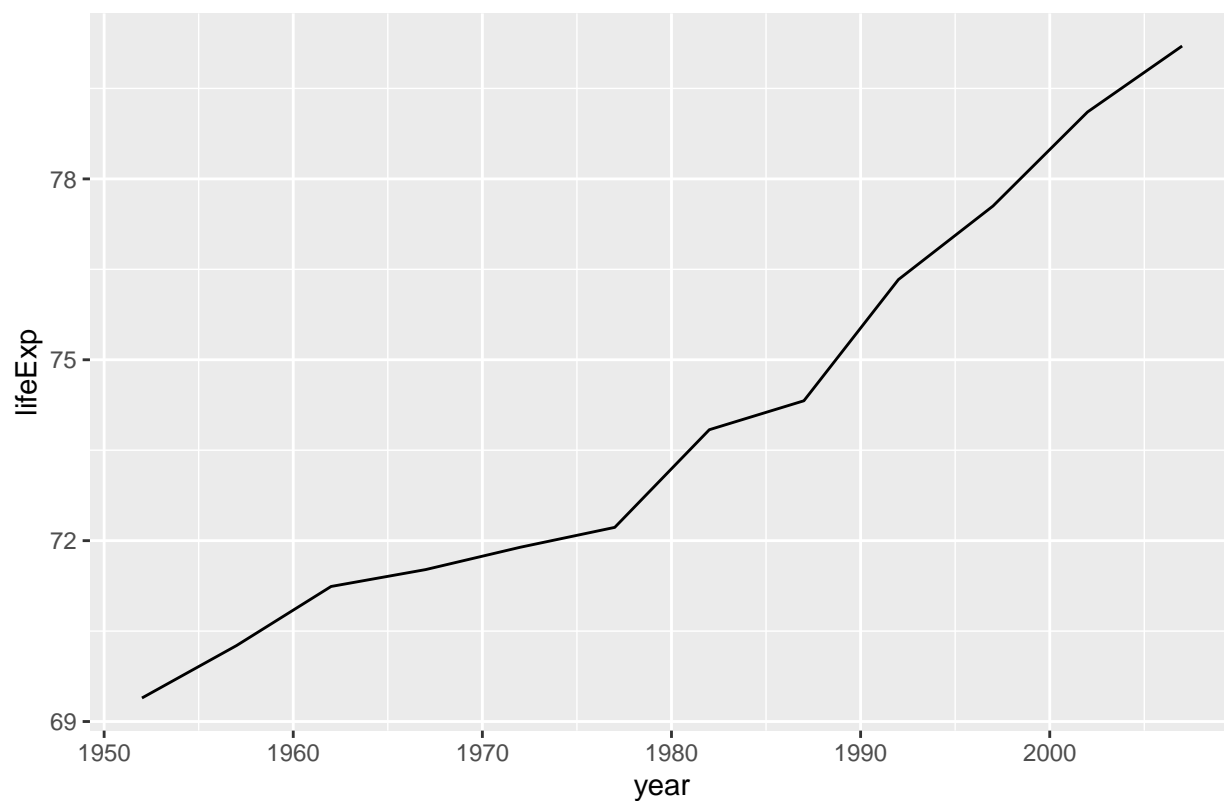
```
gapminder %>%
  ggplot(aes(year, lifeExp, group = country)) +
  geom_line(alpha = 1/3)
```



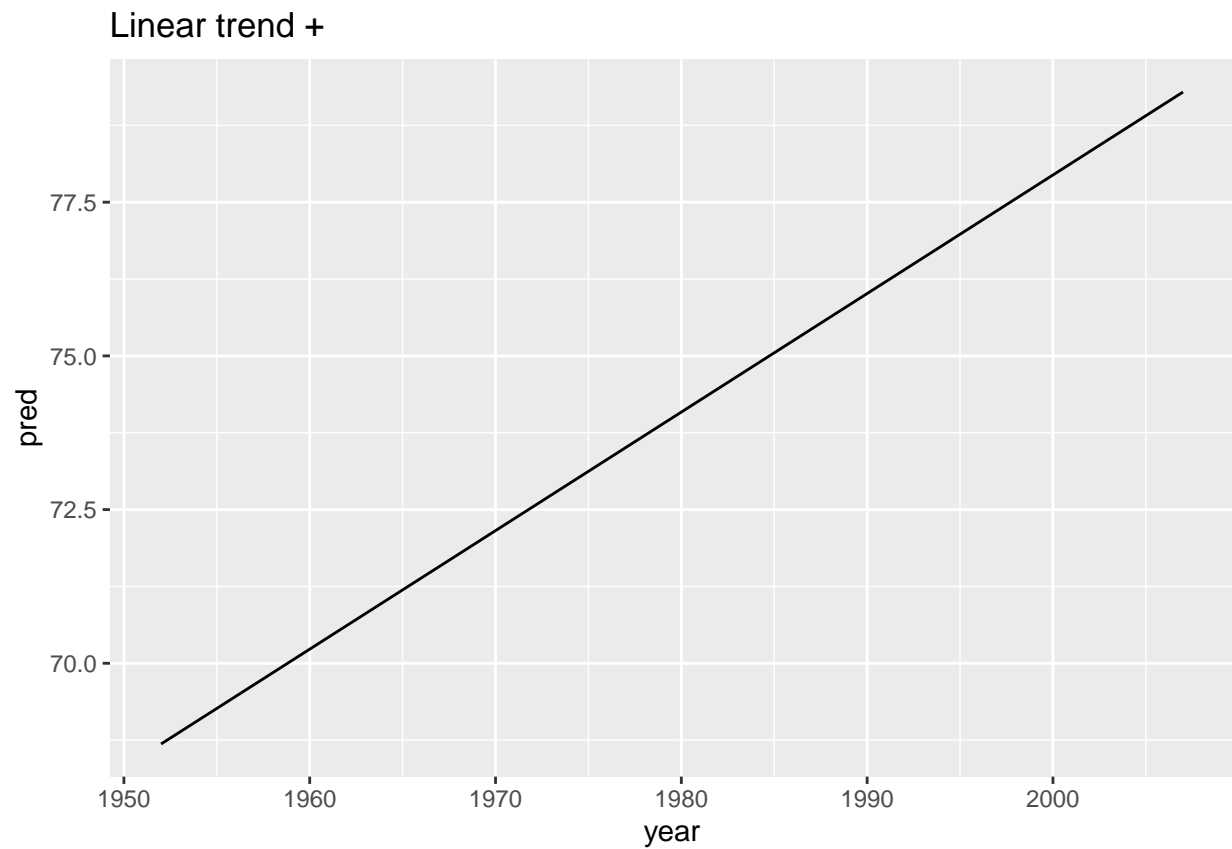
Overall we see life expectancy improving over the years, but some countries have some weird things happening. We want to explore that further. This is simple if we only have one country.

```
nz <- filter(gapminder, country == "New Zealand")
nz %>%
  ggplot(aes(year, lifeExp)) +
  geom_line() +
  ggtitle("Full data = ")
```

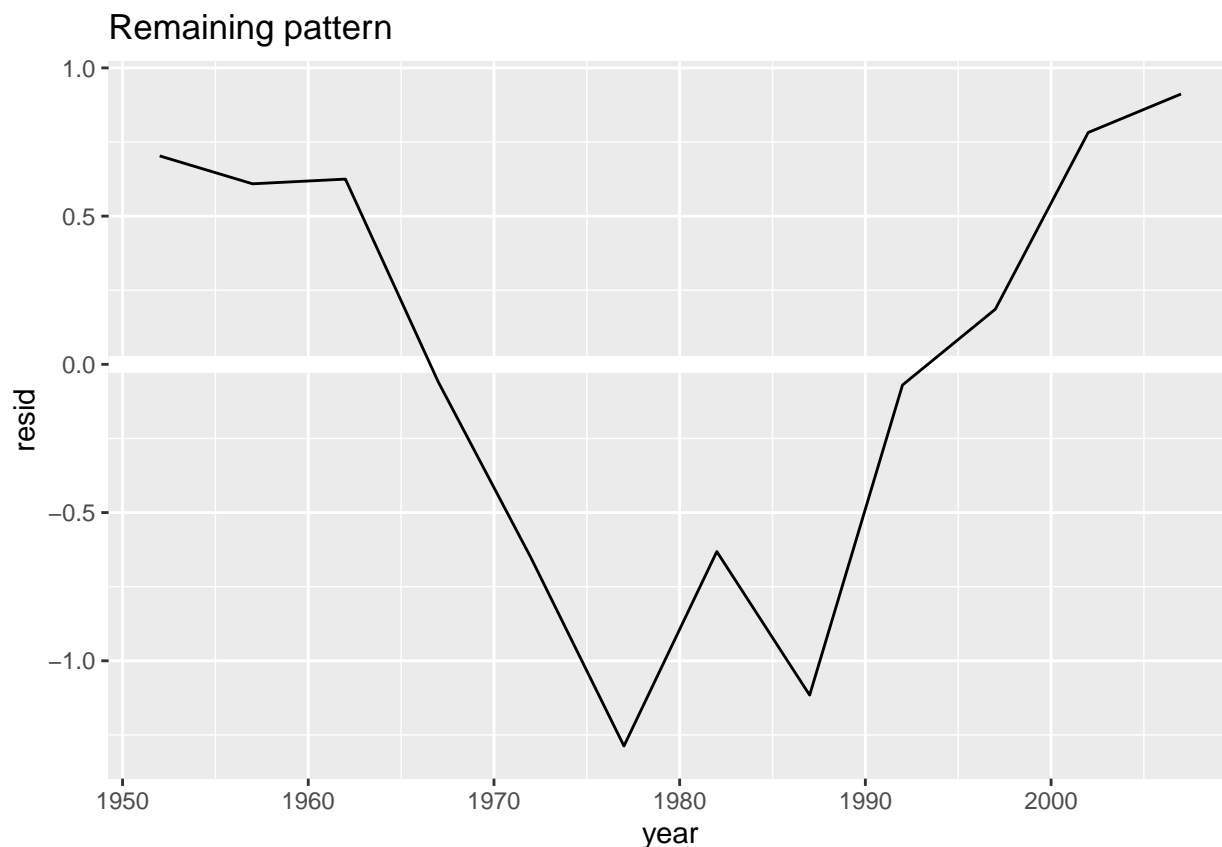
Full data =



```
nz_mod <- lm(lifeExp ~ year, data = nz)
nz %>%
  add_predictions(nz_mod) %>%
  ggplot(aes(year, pred)) +
  geom_line() +
  ggtitle("Linear trend + ")
```



```
nz %>%  
  add_residuals(nz_mod) %>%  
  ggplot(aes(year, resid)) +  
  geom_hline(yintercept = 0, colour = "white", size = 3) +  
  geom_line() +  
  ggtitle("Remaining pattern")
```



To do so for every country, we need a nested data frame.

```
by_country <- gapminder %>%
  group_by(country, continent) %>%
  nest()
```

```
by_country
```

```
## # A tibble: 142 x 3
## # Groups:   country, continent [142]
##   country    continent data
##   <fct>      <fct>    <list>
## 1 Afghanistan Asia      <tibble [12 x 4]>
## 2 Albania     Europe    <tibble [12 x 4]>
## 3 Algeria     Africa    <tibble [12 x 4]>
## 4 Angola      Africa    <tibble [12 x 4]>
## 5 Argentina   Americas  <tibble [12 x 4]>
## 6 Australia   Oceania   <tibble [12 x 4]>
## 7 Austria     Europe    <tibble [12 x 4]>
## 8 Bahrain     Asia      <tibble [12 x 4]>
## 9 Bangladesh  Asia      <tibble [12 x 4]>
## 10 Belgium    Europe    <tibble [12 x 4]>
## # ... with 132 more rows
```

We see that every country has a row, and the data column are nested data frames containing that country's observations.

We can now create a function that will create a model:

```
country_model <- function(df) {
  lm(lifeExp ~ year, data = df)
}
```

Apply it to the `by_country` data frame, and add it as an additional column.

```
models <- map(by_country$data, country_model)
by_country <- by_country %>%
  mutate(model = map(data, country_model))
by_country
```

```
## # A tibble: 142 x 4
## # Groups:   country, continent [142]
##   country    continent data          model
##   <fct>      <fct>    <list>      <list>
## 1 Afghanistan Asia    <tibble [12 x 4]> <lm>
## 2 Albania    Europe  <tibble [12 x 4]> <lm>
## 3 Algeria    Africa  <tibble [12 x 4]> <lm>
## 4 Angola     Africa  <tibble [12 x 4]> <lm>
## 5 Argentina  Americas <tibble [12 x 4]> <lm>
## 6 Australia  Oceania  <tibble [12 x 4]> <lm>
## 7 Austria    Europe  <tibble [12 x 4]> <lm>
## 8 Bahrain    Asia    <tibble [12 x 4]> <lm>
## 9 Bangladesh Asia    <tibble [12 x 4]> <lm>
## 10 Belgium   Europe  <tibble [12 x 4]> <lm>
## # ... with 132 more rows
```

Next, we can compute the residuals of each model and add it to the data frame.

```
by_country <- by_country %>%
  mutate(
    resids = map2(data, model, add_residuals)
  )
by_country
```

```
## # A tibble: 142 x 5
## # Groups:   country, continent [142]
##   country    continent data          model resids
##   <fct>      <fct>    <list>      <list> <list>
## 1 Afghanistan Asia    <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 2 Albania    Europe  <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 3 Algeria    Africa  <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 4 Angola     Africa  <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 5 Argentina  Americas <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 6 Australia  Oceania  <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 7 Austria    Europe  <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 8 Bahrain    Asia    <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 9 Bangladesh Asia    <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## 10 Belgium   Europe  <tibble [12 x 4]> <lm>    <tibble [12 x 5]>
## # ... with 132 more rows
```

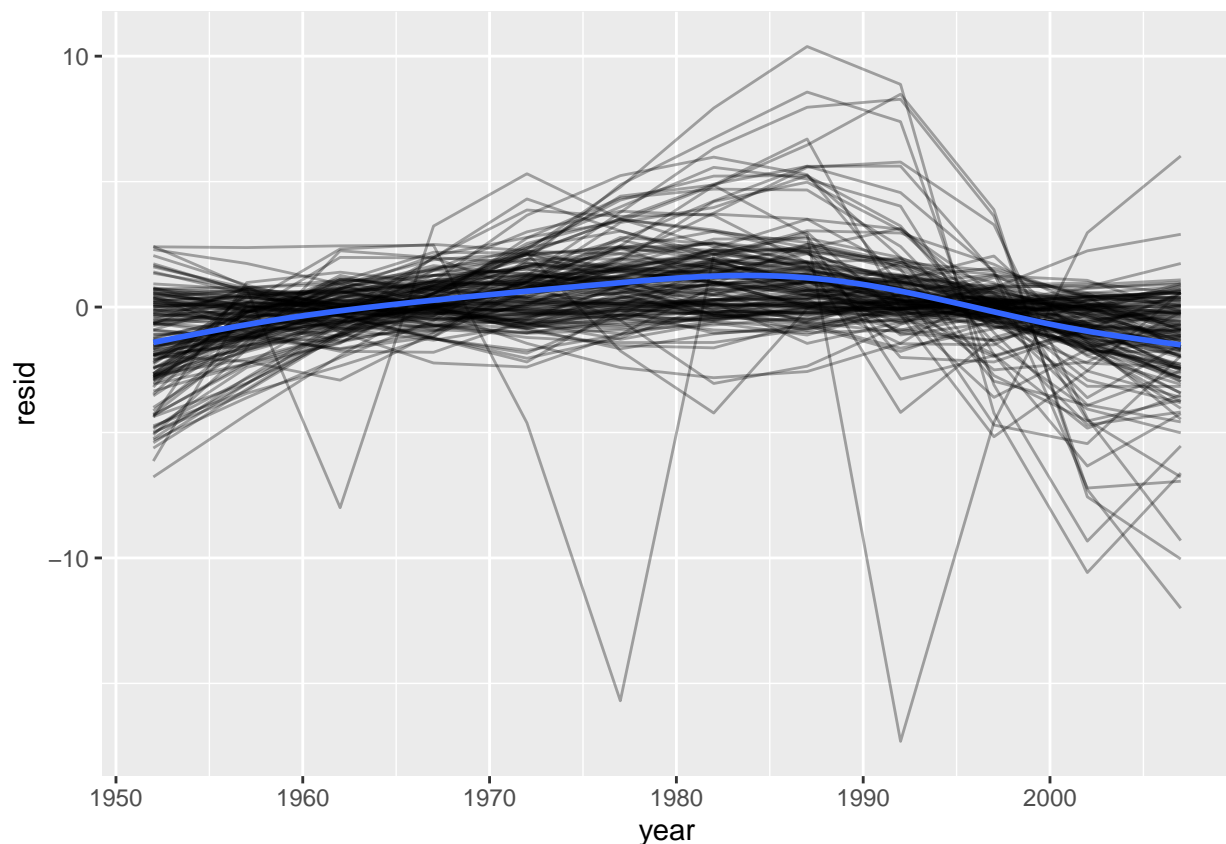
To plot the data, we need to undo the nesting with `unnest()`.

```
resids <- unnest(by_country, resids)
resids
```

```
## # A tibble: 1,704 x 9
```

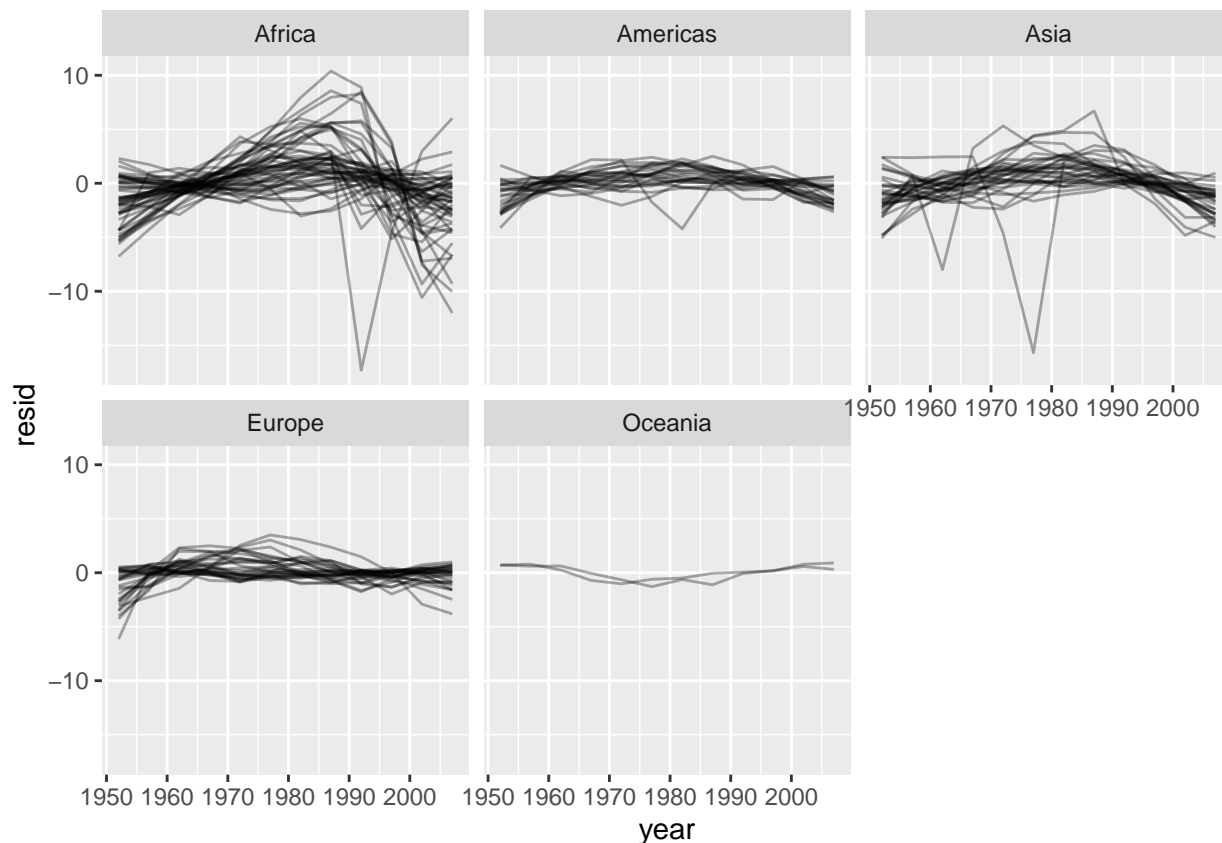
```
## # Groups:   country, continent [142]
##   country   continent data      model  year lifeExp   pop gdpPercap  resid
##   <fct>     <fct>   <list>    <list> <int>  <dbl>  <int>    <dbl>   <dbl>
## 1 Afghanis~ Asia   <tibble [1~ <lm>   1952   28.8 8.43e6    779. -1.11
## 2 Afghanis~ Asia   <tibble [1~ <lm>   1957   30.3 9.24e6    821. -0.952
## 3 Afghanis~ Asia   <tibble [1~ <lm>   1962   32.0 1.03e7    853. -0.664
## 4 Afghanis~ Asia   <tibble [1~ <lm>   1967   34.0 1.15e7    836. -0.0172
## 5 Afghanis~ Asia   <tibble [1~ <lm>   1972   36.1 1.31e7    740.  0.674
## 6 Afghanis~ Asia   <tibble [1~ <lm>   1977   38.4 1.49e7    786.  1.65
## 7 Afghanis~ Asia   <tibble [1~ <lm>   1982   39.9 1.29e7    978.  1.69
## 8 Afghanis~ Asia   <tibble [1~ <lm>   1987   40.8 1.39e7    852.  1.28
## 9 Afghanis~ Asia   <tibble [1~ <lm>   1992   41.7 1.63e7    649.  0.754
## 10 Afghanis~ Asia  <tibble [1~ <lm>   1997   41.8 2.22e7    635. -0.534
## # ... with 1,694 more rows
```

```
resids %>%
  ggplot(aes(year, resid)) +
  geom_line(aes(group = country), alpha = 1 / 3) +
  geom_smooth(se = FALSE)
```



We can split this graph up by continents.

```
resids %>%
  ggplot(aes(year, resid, group = country)) +
  geom_line(alpha = 1 / 3) +
  facet_wrap(~continent)
```



Unsurprisingly, we see large residuals in Africa and Asia, but small residuals in the Americas and Europe. Not much can be said about Oceania, as there are only two countries. To get information on the quality of the models, we can use the broom package:

```
glance <- by_country %>%
  mutate(glance = map(model, broom::glance)) %>%
  unnest(glance, .drop = TRUE)
glance
```

```
## # A tibble: 142 x 17
## # Groups:   country, continent [142]
##   country continent data model resid r.squared adj.r.squared sigma statistic
##   <fct> <fct> <lis> <lis> <list> <dbl> <dbl> <dbl> <dbl>
## 1 Afghan~ Asia <tib~ <lm> <tibb~ 0.948 0.942 1.22 181.
## 2 Albania Europe <tib~ <lm> <tibb~ 0.911 0.902 1.98 102.
## 3 Algeria Africa <tib~ <lm> <tibb~ 0.985 0.984 1.32 662.
## 4 Angola Africa <tib~ <lm> <tibb~ 0.888 0.877 1.41 79.1
## 5 Argent~ Americas <tib~ <lm> <tibb~ 0.996 0.995 0.292 2246.
## 6 Austra~ Oceania <tib~ <lm> <tibb~ 0.980 0.978 0.621 481.
## 7 Austria Europe <tib~ <lm> <tibb~ 0.992 0.991 0.407 1261.
## 8 Bahrain Asia <tib~ <lm> <tibb~ 0.967 0.963 1.64 291.
## 9 Bangla~ Asia <tib~ <lm> <tibb~ 0.989 0.988 0.977 930.
## 10 Belgium Europe <tib~ <lm> <tibb~ 0.995 0.994 0.293 1822.
## # ... with 132 more rows, and 8 more variables: p.value <dbl>, df <dbl>,
## # logLik <dbl>, AIC <dbl>, BIC <dbl>, deviance <dbl>, df.residual <int>,
## # nob <int>
```

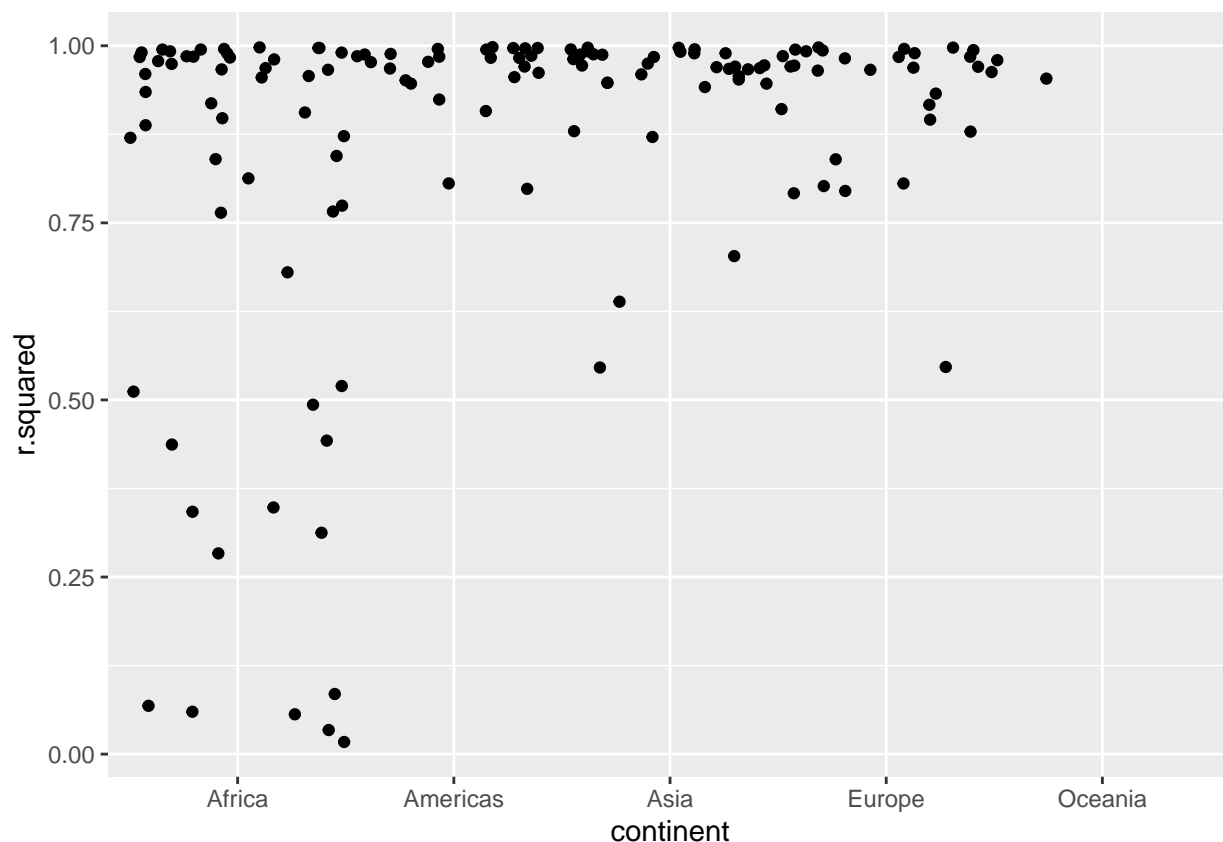
R-squared is a goodness-of-fit measure, so to see the poorest quality models, we can arrange by R-squared.


```
glance %>%
  arrange(r.squared)
```

```
## # A tibble: 142 x 17
## # Groups:   country, continent [142]
##   country continent data model resid r.squared adj.r.squared sigma statistic
##   <fct>   <fct>   <lis> <lis> <list>   <dbl>         <dbl> <dbl>   <dbl>
## 1 Rwanda  Africa   <tib~ <lm>  <tibb~    0.0172      -0.0811  6.56    0.175
## 2 Botswa~ Africa   <tib~ <lm>  <tibb~    0.0340      -0.0626  6.11    0.352
## 3 Zimbab~ Africa   <tib~ <lm>  <tibb~    0.0562      -0.0381  7.21    0.596
## 4 Zambia  Africa   <tib~ <lm>  <tibb~    0.0598      -0.0342  4.53    0.636
## 5 Swazil~ Africa   <tib~ <lm>  <tibb~    0.0682      -0.0250  6.64    0.732
## 6 Lesotho Africa   <tib~ <lm>  <tibb~    0.0849      -0.00666 5.93    0.927
## 7 Cote d~ Africa   <tib~ <lm>  <tibb~    0.283        0.212    3.93    3.95
## 8 South ~ Africa   <tib~ <lm>  <tibb~    0.312        0.244    4.74    4.54
## 9 Uganda  Africa   <tib~ <lm>  <tibb~    0.342        0.276    3.19    5.20
## 10 Congo,~ Africa   <tib~ <lm>  <tibb~    0.348        0.283    2.43    5.34
## # ... with 132 more rows, and 8 more variables: p.value <dbl>, df <dbl>,
## #   logLik <dbl>, AIC <dbl>, BIC <dbl>, deviance <dbl>, df.residual <int>,
## #   nobs <int>
```

All the worst models are in Africa, so lets plot and see:

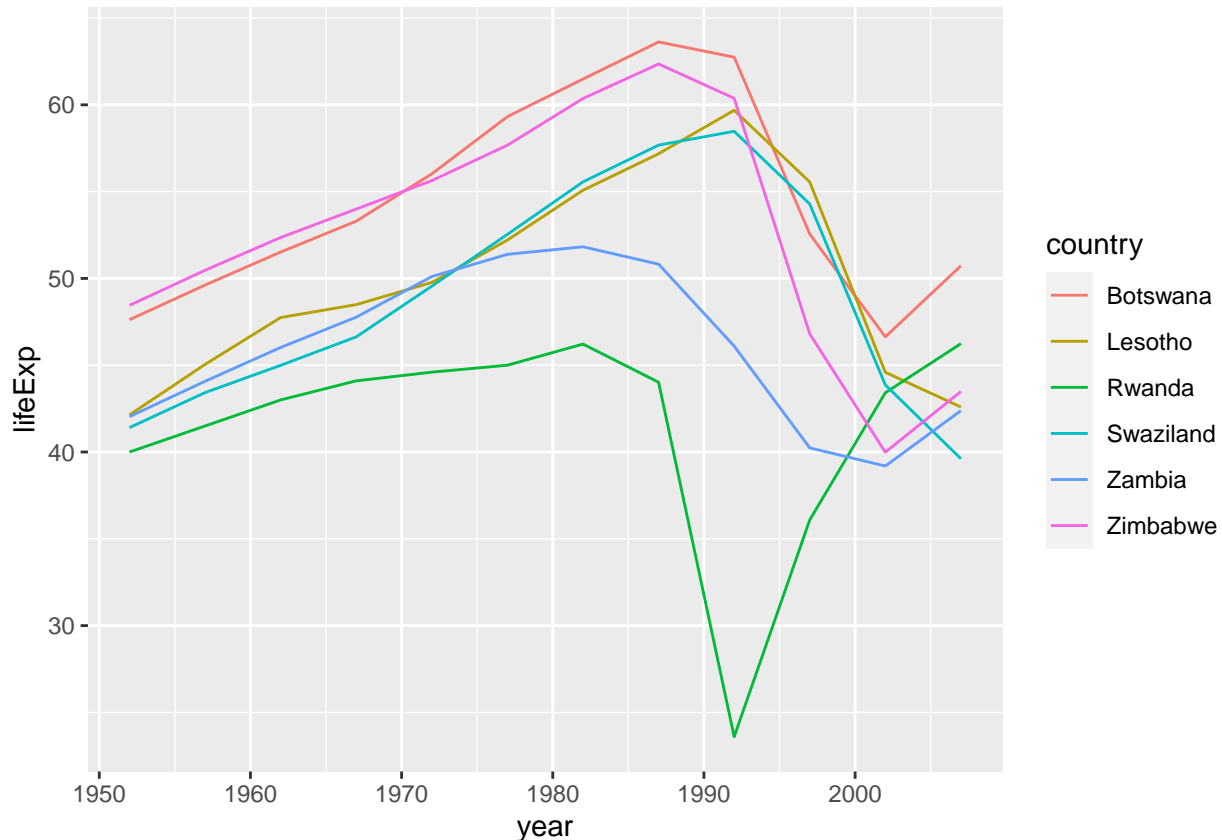
```
glance %>%
  ggplot(aes(continent, r.squared)) +
  geom_jitter(width = 0.5)
```



Finally, lets pull out the worst models and plot them:

```
bad_fit <- filter(glance, r.squared < 0.25)

gapminder %>%
  semi_join(bad_fit, by = "country") %>%
  ggplot(aes(year, lifeExp, colour = country)) +
  geom_line()
```



The HIV/AIDS epidemic and the Rwandan genocide explain a lot of these trends.

Exercise 2 (Website: 25.2.5 Ex. 3)

To create the last plot in section 25.2 gapminder (showing the data for the countries with the worst model fits according to R^2), we needed two steps; we created a data frame with one row per country and then semi-joined it to the original dataset. It's possible to avoid using this join if we use `unnest()` instead of `unnest(.drop = TRUE)` when creating the `glance` tibble. Demonstrate how this could be done.

(Hint: You might have to use another `unnest()` call. Also note that you won't need the `bad_fit` tibble, since you will be able to pipe directly into the `filter()`.)

```
gapminder %>%
  group_by(country, continent) %>%
  nest() %>%
  mutate(model = map(data, ~lm(lifeExp ~ year, .))) %>%
  mutate(glance = map(model, broom::glance)) %>%
  unnest(data) %>%
  unnest(glance) %>%
  filter(r.squared < 0.25) %>%
```

```
ggplot(mapping = aes(x = year, y = lifeExp)) +  
geom_line(mapping = aes(color = country))
```

